

MIT Art Design and Technology University
MIT School of Computing, Pune

**Department of Computer Science and
Engineering**

Lab Report

Course- HPC Lab

Class - L.Y. (SEM-I), Core, AIEC

Name: Sourav Shailesh Toshniwal

Roll No: 2213047

A.Y. 2024 - 2025

Lab Experiment List

Sr. No.	Name of Experiment	CO
1.	Familiarization with Linux commands	CO1
2.	Familiarization with SLURM commands	CO1
3.	Write an OpenMP program to print Hello world with thread ID.	CO2
4.	Write your first Parallel Program, with which you should be able to print your NAME from 4 underline cores.	CO2
5.	<p>Write a C program utilizing OpenMP directives to demonstrate the behavior of the private clause.</p> <ul style="list-style-type: none"> ▪ The program should perform the following steps: ▪ Initialize OpenMP with 4 threads. ▪ Declare an integer variable val and initialize it to a value of 1234. ▪ Print the initial value of val outside the OpenMP parallel region. ▪ Enter an OpenMP parallel region using the omp parallel directive, with the firstprivate clause applied to the variable val. ▪ Inside the parallel region, each thread should print the current value of val, increment it by 1, and then print the updated value. ▪ Print the final value of val outside the parallel region. 	CO2
6.	<p>Write a C program utilizing OpenMP directives to demonstrate the behavior of the private clause.</p> <p>Steps to follow :</p> <ul style="list-style-type: none"> ▪ Open text editor. ▪ write the below program in it. ▪ Save the file with .c extention. 	CO2

	<ul style="list-style-type: none"> ▪ Compile and execute with given commands. 	
7.	Write a Parallel C program where the iterations of a loop should be scheduled statically across the team of threads. A thread should perform CHUNK iterations at a time before being scheduled for the next CHUNK of work.	CO2
8.	Write a Parallel C program which should print the series of 2 and 4. Make sure both should be executed by different threads.	CO2
9.	Write MPI Program to print "Hello World". MPI program to send and receive Hello World messages from all other processes to a Root process and print the received messages.	CO3
10.	MPI program to send two numbers (array elements) per process to a Root process and print the received messages.	CO3
11.	MPI program to find sum of first N integers using any given number of processes. Example, N=10,000 and no. of processes can be 4 or 8 or 12 etc.	CO4
12.	MPI program to find sum of n integers on in which processors are arranged in ring topology using MPI point-to-point blocking communication library calls.	CO4
13.	Write a CUDA program to perform two matrix addition.	CO5
14.	Write a CUDA program to perform two matrix multiplication.	CO5

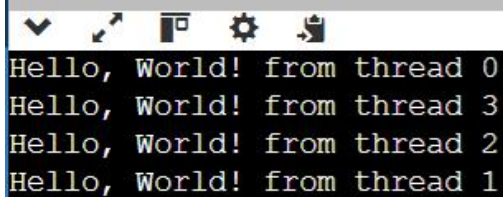
Experiment No 3

Experiment Title: OpenMP program 1

Problem Statement: Write an OpenMP program to print Hello world with thread ID.

Source Code and Output /Screenshots:

```
1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Start parallel region
6      #pragma omp parallel
7      {
8          // Get the thread ID of the current thread
9          int thread_id = omp_get_thread_num();
10
11         // Print Hello World along with the thread ID
12         printf("Hello, World! from thread %d\n", thread_id);
13     }
14
15     return 0;
16 }
```



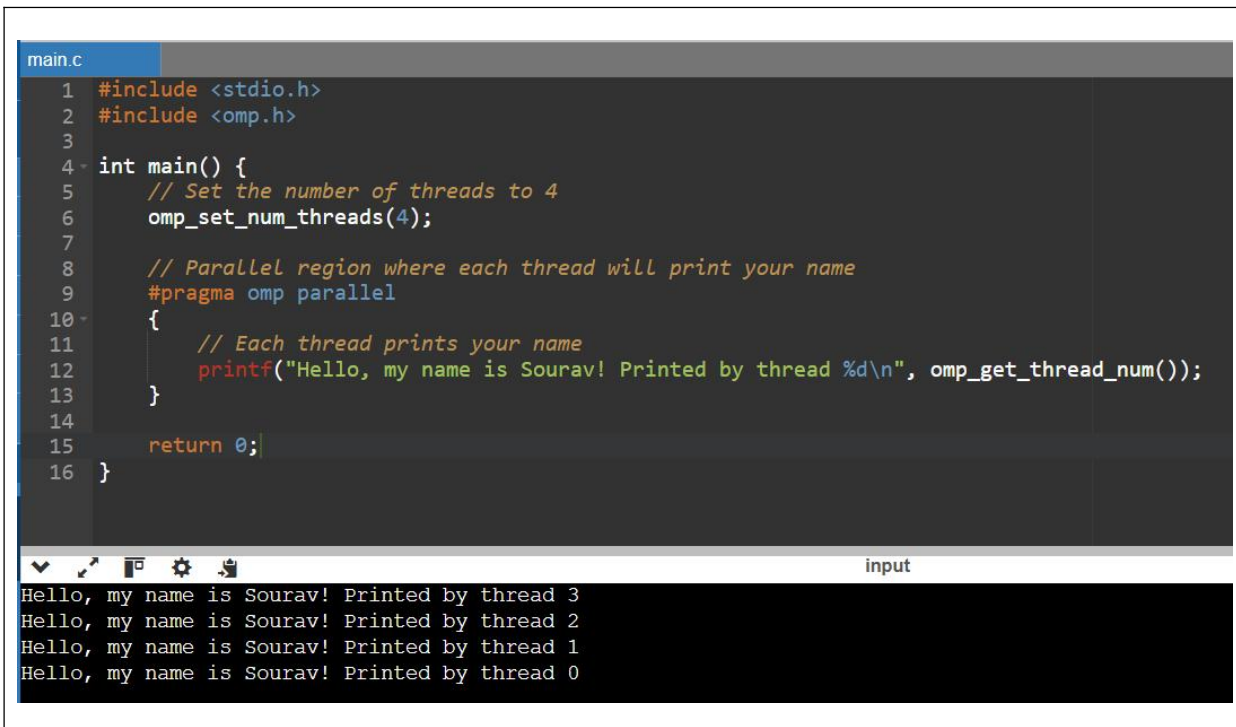
```
Hello, World! from thread 0
Hello, World! from thread 3
Hello, World! from thread 2
Hello, World! from thread 1
```

Experiment No 4

Experiment Title: OpenMP Program 2

Problem Statement: Write your first Parallel Program, with which you should be able to print your NAME from 4 underline cores.

Source Code and Output /Screenshots:



The screenshot displays a code editor window titled 'main.c' with a dark background. The code is written in C and uses OpenMP for parallelization. It includes `<stdio.h>` and `<omp.h>`. The `main` function sets the number of threads to 4 using `omp_set_num_threads(4);`. A parallel region is defined with `#pragma omp parallel`, containing a block where each thread prints a message: `printf("Hello, my name is Sourav! Printed by thread %d\n", omp_get_thread_num());`. The program returns 0. Below the code editor, the output is shown in a terminal window, displaying four lines of the same message, each corresponding to a different thread ID (3, 2, 1, and 0).

```
main.c
1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Set the number of threads to 4
6      omp_set_num_threads(4);
7
8      // Parallel region where each thread will print your name
9      #pragma omp parallel
10     {
11         // Each thread prints your name
12         printf("Hello, my name is Sourav! Printed by thread %d\n", omp_get_thread_num());
13     }
14
15     return 0;
16 }
```

input

```
Hello, my name is Sourav! Printed by thread 3
Hello, my name is Sourav! Printed by thread 2
Hello, my name is Sourav! Printed by thread 1
Hello, my name is Sourav! Printed by thread 0
```

Experiment No 5

Experiment Title: OpenMP Program 3

Problem Statement:

Write a C program utilizing OpenMP directives to demonstrate the behavior of the private clause. The program should perform the following steps:

- Initialize OpenMP with 4 threads.
- Declare an integer variable val and initialize it to a value of 1234.
- Print the initial value of val outside the OpenMP parallel region.
- Enter an OpenMP parallel region using the omp parallel directive, with the firstprivate clause applied to the variable val.
- Inside the parallel region, each thread should print the current value of val, increment it by 1, and then print the updated value.
- Print the final value of val outside the parallel region.

Source Code and Output /Screenshots:

```

1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Initialize OpenMP with 4 threads
6      omp_set_num_threads(4);
7
8      // Declare an integer variable val and initialize it to 1234
9      int val = 1234;
10
11     // Print the initial value of val outside the parallel region
12     printf("Initial value of val (outside parallel region): %d\n", val);
13
14     // Enter an OpenMP parallel region with the firstprivate clause applied to val
15     #pragma omp parallel firstprivate(val)
16     {
17         // Get the thread number
18         int thread_id = omp_get_thread_num();
19
20         // Print the current value of val (private copy for each thread)
21         printf("Thread %d: Current value of val = %d\n", thread_id, val);
22
23         // Increment val by 1 (this only affects the private copy for this thread)
24         val++;
25
26         // Print the updated value of val (private copy after increment)
27         printf("Thread %d: Updated value of val = %d\n", thread_id, val);
28     }
29
30     // Print the final value of val outside the parallel region
31     printf("Final value of val (outside parallel region): %d\n", val);
32
33     return 0;
34 }

```

```
Initial value of val (outside parallel region): 1234
Thread 2: Current value of val = 1234
Thread 2: Updated value of val = 1235
Thread 0: Current value of val = 1234
Thread 0: Updated value of val = 1235
Thread 1: Current value of val = 1234
Thread 1: Updated value of val = 1235
Thread 3: Current value of val = 1234
Thread 3: Updated value of val = 1235
Final value of val (outside parallel region): 1234
```


Experiment No 6

Experiment Title: OpenMP Program 4

Problem Statement:

Write a C program utilizing OpenMP directives to demonstrate the behavior of the private clause.

Steps to follow :

Open text editor.

write the below program in it.

Save the file with .c extention.

Compile and execute with given commands.

Source Code and Output /Screenshots:

```
1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Initialize OpenMP with 4 threads
6      omp_set_num_threads(4);
7
8      // Declare an integer variable val and initialize it to 1234
9      int val = 1234;
10
11     // Print the initial value of val outside the parallel region
12     printf("Initial value of val (outside parallel region): %d\n", val);
13
14     // Enter an OpenMP parallel region with the private clause applied to val
15     #pragma omp parallel private(val)
16     {
17         // Get the thread number
18         int thread_id = omp_get_thread_num();
19
20         // Set the private copy of val in each thread (no initial value)
21         val = thread_id * 10; // Each thread sets its private val to something unique
22
23         // Print the current value of val (private copy for each thread)
24         printf("Thread %d: Current value of val = %d\n", thread_id, val);
25     }
26
27     // Print the final value of val outside the parallel region
28     printf("Final value of val (outside parallel region): %d\n", val);
29
30     return 0;
31 }
```

```
Initial value of val (outside parallel region): 1234
Thread 1: Current value of val = 10
Thread 3: Current value of val = 30
Thread 0: Current value of val = 0
Thread 2: Current value of val = 20
Final value of val (outside parallel region): 1234
```


Experiment No 7

Experiment Title: OpenMP Program 5

Problem Statement: Write a Parallel C program where the iterations of a loop should be scheduled statically across the team of threads. A thread should perform CHUNK iterations at a time before being scheduled for the next CHUNK of work.

Source Code and Output /Screenshots:

```
1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Define the number of iterations in the loop
6      int num_iterations = 100;
7
8      // Define the chunk size
9      int chunk = 10;
10
11     // Initialize OpenMP with a specific number of threads (e.g., 4 threads)
12     omp_set_num_threads(4);
13
14     // Print a message indicating the start of parallel execution
15     printf("Starting parallel loop with static scheduling and chunk size %d:\n", chunk);
16
17     // Parallel for loop with static scheduling and chunk size
18     #pragma omp parallel for schedule(static, chunk)
19     for (int i = 0; i < num_iterations; i++) {
20         // Get the thread ID
21         int thread_id = omp_get_thread_num();
22
23         // Print the iteration index and the corresponding thread ID
24         printf("Thread %d: Iteration %d\n", thread_id, i);
25     }
26
27     // Print a message indicating the end of the parallel execution
28     printf("Parallel loop execution completed.\n");
29
30     return 0;
31 }
```

```
Starting parallel loop with static scheduling and chunk size 10:
```

```
Thread 3: Iteration 30  
Thread 3: Iteration 31  
Thread 3: Iteration 32  
Thread 3: Iteration 33  
Thread 3: Iteration 34  
Thread 3: Iteration 35  
Thread 3: Iteration 36  
Thread 3: Iteration 37  
Thread 3: Iteration 38  
Thread 3: Iteration 39  
Thread 3: Iteration 70  
Thread 3: Iteration 71  
Thread 3: Iteration 72  
Thread 3: Iteration 73  
Thread 3: Iteration 74  
Thread 3: Iteration 75  
Thread 3: Iteration 76  
Thread 3: Iteration 77  
Thread 3: Iteration 78  
Thread 3: Iteration 79  
Thread 1: Iteration 10  
Thread 1: Iteration 11  
Thread 1: Iteration 12  
Thread 1: Iteration 13  
Thread 1: Iteration 14  
Thread 1: Iteration 15  
Thread 1: Iteration 16  
Thread 1: Iteration 17  
Thread 1: Iteration 18  
Thread 1: Iteration 19  
Thread 1: Iteration 50  
Thread 1: Iteration 51  
Thread 1: Iteration 52  
Thread 1: Iteration 53  
Thread 1: Iteration 54
```

```
Thread 1: Iteration 55
Thread 1: Iteration 56
Thread 1: Iteration 57
Thread 1: Iteration 58
Thread 1: Iteration 59
Thread 1: Iteration 90
Thread 1: Iteration 91
Thread 1: Iteration 92
Thread 1: Iteration 93
Thread 1: Iteration 94
Thread 1: Iteration 95
Thread 1: Iteration 96
Thread 1: Iteration 97
Thread 1: Iteration 98
Thread 1: Iteration 99
Thread 2: Iteration 20
Thread 2: Iteration 21
Thread 2: Iteration 22
Thread 2: Iteration 23
Thread 2: Iteration 24
Thread 2: Iteration 25
Thread 2: Iteration 26
Thread 2: Iteration 27
Thread 2: Iteration 28
Thread 2: Iteration 29
Thread 2: Iteration 60
Thread 2: Iteration 61
Thread 2: Iteration 62
Thread 2: Iteration 63
Thread 2: Iteration 64
Thread 2: Iteration 65
Thread 2: Iteration 66
Thread 2: Iteration 67
Thread 2: Iteration 68
Thread 2: Iteration 69
Thread 0: Iteration 0
```

```
Thread 0: Iteration 1  
Thread 0: Iteration 2  
Thread 0: Iteration 3  
Thread 0: Iteration 4  
Thread 0: Iteration 5  
Thread 0: Iteration 6  
Thread 0: Iteration 7  
Thread 0: Iteration 8  
Thread 0: Iteration 9  
Thread 0: Iteration 40  
Thread 0: Iteration 41  
Thread 0: Iteration 42  
Thread 0: Iteration 43  
Thread 0: Iteration 44  
Thread 0: Iteration 45  
Thread 0: Iteration 46  
Thread 0: Iteration 47  
Thread 0: Iteration 48  
Thread 0: Iteration 49  
Thread 0: Iteration 80  
Thread 0: Iteration 81  
Thread 0: Iteration 82  
Thread 0: Iteration 83  
Thread 0: Iteration 84  
Thread 0: Iteration 85  
Thread 0: Iteration 86  
Thread 0: Iteration 87  
Thread 0: Iteration 88  
Thread 0: Iteration 89  
Parallel loop execution completed.
```

Experiment No 8

Experiment Title: OpenMP Program 6

Problem Statement:

Write a Parallel C program which should print the series of 2 and 4. Make sure both should be executed by different threads.

Source Code and Output /Screenshots:

```
1  #include <stdio.h>
2  #include <omp.h>
3
4  int main() {
5      // Set the number of threads to 2 (one for each series)
6      omp_set_num_threads(2);
7
8      // Parallel region where each thread will execute different tasks
9      #pragma omp parallel
10     {
11         int thread_id = omp_get_thread_num();
12
13         // If thread_id is 0, print the series of 2
14         if (thread_id == 0) {
15             for (int i = 1; i <= 5; i++) { // Print 5 terms of the series 2, 2, 2, ...
16                 printf("Thread %d: %d\n", thread_id, 2);
17             }
18         }
19         // If thread_id is 1, print the series of 4
20         else if (thread_id == 1) {
21             for (int i = 1; i <= 5; i++) { // Print 5 terms of the series 4, 4, 4, ...
22                 printf("Thread %d: %d\n", thread_id, 4);
23             }
24         }
25     }
26
27     return 0;
28 }
```

```
Thread 1: 4
Thread 1: 4
Thread 1: 4
Thread 1: 4
Thread 1: 4
Thread 0: 2
Thread 0: 2
Thread 0: 2
Thread 0: 2
Thread 0: 2
```

Experiment No 9

Experiment Title: MPI Program 1

Problem Statement: Write MPI Program to print "Hello World".

MPI program to send and receive Hello World messages from all other processes to a Root process and print the received messages.

Source Code and Output /Screenshots:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    int rank, size;
    char message[20];
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    sprintf(message, "Hello from process %d", rank);

    if (rank == 0) {
        printf("Root process %d: Receiving messages:\n", rank);
        for (int i = 1; i < size; i++) {
            MPI_Recv(message, 20, MPI_CHAR, i, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("Received message: %s\n", message);
        }
        printf("Received message: Hello from process 0\n");
    }
    else {
        MPI_Send(message, 20, MPI_CHAR, 0, 0, MPI_COMM_WORLD);
    }
    MPI_Finalize();

    return 0;
}
```

```
Root process 0: Receiving messages:
Received message: Hello from process 1
Received message: Hello from process 2
Received message: Hello from process 3
Received message: Hello from process 0
```


Experiment No 10

Experiment Title: MPI Program 2

Problem Statement: MPI program to send two numbers (array elements) per process to a Root process and print the received messages.

Source Code and Output /Screenshots:

```
#include <mpi.h>
#include <stdio.h>
#define NUM_ELEMENTS 2

int main(int argc, char** argv) {
    int rank, size;
    int data[NUM_ELEMENTS];
    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    data[0] = rank * 10 + 1;
    data[1] = rank * 10 + 2;

    if (rank == 0) {
        printf("Root process %d: Receiving messages:\n", rank);
        for (int i = 1; i < size; i++) {
            MPI_Recv(data, NUM_ELEMENTS, MPI_INT, i, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
            printf("Received from process %d: %d, %d\n", i, data[0], data[1]);
        }
        printf("Received from process %d: %d, %d\n", rank, data[0], data[1]);
    }
    else {
        MPI_Send(data, NUM_ELEMENTS, MPI_INT, 0, 0, MPI_COMM_WORLD);
    }

    MPI_Finalize();

    return 0;
}
```

Root process 0: Receiving messages:

Received from process 1: 11, 12

Received from process 2: 21, 22

Received from process 3: 31, 32

Received from process 0: 1, 2

Experiment No 11

Experiment Title: MPI Program 3

Problem Statement: MPI program to find sum of first N integers using any given number of processes. Example, N=10,000 and no. of processes can be 4 or 8 or 12 etc.

Source Code and Output /Screenshots:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    int rank, size;
    long long N = 10000;
    long long local_sum = 0, global_sum = 0;
    long long start, end;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    long long range_per_process = N / size;
    start = rank * range_per_process + 1;
    end = (rank + 1) * range_per_process;
    if (rank == size - 1) {
        end = N;
    }
    for (long long i = start; i <= end; i++) {
        local_sum += i;
    }
    MPI_Reduce(&local_sum, &global_sum, 1, MPI_LONG_LONG, MPI_SUM, 0, MPI_COMM_WORLD);

    if (rank == 0) {
        printf("The sum of the first %lld integers is: %lld\n", N, global_sum);
    }
    MPI_Finalize();

    return 0;
}
```

The sum of the first 10000 integers is: 50005000

Experiment No 12

Experiment Title: MPI Program 4

Problem Statement: MPI program to find sum of n integers in which processors are arranged in ring topology using MPI point-to-point blocking communication library calls.

Source Code and Output /Screenshots:

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    int rank, size;
    long long N = 10000;
    long long local_sum = 0, total_sum = 0;
    long long start, end;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    long long range_per_process = N / size;
    start = rank * range_per_process + 1;
    end = (rank + 1) * range_per_process;
    if (rank == size - 1) {
        end = N;
    }
    for (long long i = start; i <= end; i++) {
        local_sum += i;
    }
    int prev = (rank - 1 + size) % size;
    int next = (rank + 1) % size;
    if (rank == 0) {
        MPI_Send(&local_sum, 1, MPI_LONG_LONG, next, 0, MPI_COMM_WORLD);
        MPI_Recv(&total_sum, 1, MPI_LONG_LONG, prev, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        total_sum += local_sum;
        printf("The total sum of the first %lld integers is: %lld\n", N, total_sum);
    }
    else {
        MPI_Send(&local_sum, 1, MPI_LONG_LONG, next, 0, MPI_COMM_WORLD);
        MPI_Recv(&total_sum, 1, MPI_LONG_LONG, prev, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        total_sum += local_sum;
        if (rank != size - 1) {
            MPI_Send(&total_sum, 1, MPI_LONG_LONG, next, 0, MPI_COMM_WORLD);
        }
    }
    MPI_Finalize();

    return 0;
}
```

The total sum of the first 10000 integers is: 50005000

Experiment No 13

Experiment Title: CUDA Program 1

Problem Statement: Write a CUDA program to perform two matrix additions.

Source Code and Output /Screenshots:

```
#include <stdio.h>
#include <cuda_runtime.h>
#define N 1024
__global__ void matrixAdd(float *A, float *B, float *C, int width) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    int idy = blockIdx.y * blockDim.y + threadIdx.y;

    if (idx < width && idy < width) {
        int index = idy * width + idx;
        C[index] = A[index] + B[index];
    }
}

int main() {
    int size = N * N * sizeof(float);
    float *h_A, *h_B, *h_C, *h_D, *h_E;
    h_A = (float*)malloc(size);
    h_B = (float*)malloc(size);
    h_C = (float*)malloc(size);
    h_D = (float*)malloc(size);
    h_E = (float*)malloc(size);
    for (int i = 0; i < N * N; i++) {
        h_A[i] = rand() % 10;
        h_B[i] = rand() % 10;
        h_D[i] = rand() % 10;
    }
    float *d_A, *d_B, *d_C, *d_D, *d_E;
    cudaMalloc(&d_A, size);
    cudaMalloc(&d_B, size);
    cudaMalloc(&d_C, size);
    cudaMalloc(&d_D, size);
    cudaMalloc(&d_E, size);
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
```

```

cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_D, h_D, size, cudaMemcpyHostToDevice);
dim3 threadsPerBlock(16, 16);
dim3 numBlocks((N + 15) / 16, (N + 15) / 16);
matrixAdd<<<numBlocks, threadsPerBlock>>>(d_A, d_B, d_C, N);

cudaDeviceSynchronize();
matrixAdd<<<numBlocks, threadsPerBlock>>>(d_C, d_D, d_E, N);
cudaDeviceSynchronize();
cudaMemcpy(h_E, d_E, size, cudaMemcpyDeviceToHost);
printf("Result matrix E (a few elements):\n");
for (int i = 0; i < 5; i++) {
    printf("%f ", h_E[i]);
}
printf("\n");
free(h_A);
free(h_B);
free(h_C);
free(h_D);
free(h_E);
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);
cudaFree(d_D);
cudaFree(d_E);

return 0;
}

```

Result matrix E (a few elements):

42.000000 32.000000 24.000000 55.000000 19.000000

Experiment No 14

Experiment Title: CUDA Program 2

Problem Statement: Write a CUDA program to perform two matrix multiplications.

Source Code and Output /Screenshots:

```
#include <stdio.h>
#include <cuda_runtime.h>

#define N 1024
__global__ void matrixMultiply(float *A, float *B, float *C, int width) {
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < width && col < width) {
        float value = 0;
        for (int k = 0; k < width; k++) {
            value += A[row * width + k] * B[k * width + col];
        }
        C[row * width + col] = value;
    }
}

int main() {
    int size = N * N * sizeof(float);
    float *h_A, *h_B, *h_C, *h_D, *h_E;
    h_A = (float*)malloc(size);
    h_B = (float*)malloc(size);
    h_C = (float*)malloc(size);
    h_D = (float*)malloc(size);
    h_E = (float*)malloc(size);
    for (int i = 0; i < N * N; i++) {
        h_A[i] = rand() % 10;
        h_B[i] = rand() % 10;
        h_D[i] = rand() % 10;
    }
    float *d_A, *d_B, *d_C, *d_D, *d_E;
    cudaMalloc(&d_A, size);
    cudaMalloc(&d_B, size);
    cudaMalloc(&d_C, size);
```

```

cudaMalloc(&d_A, size);
cudaMalloc(&d_B, size);
cudaMalloc(&d_C, size);
cudaMalloc(&d_D, size);
cudaMalloc(&d_E, size);
cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
cudaMemcpy(d_D, h_D, size, cudaMemcpyHostToDevice);
dim3 threadsPerBlock(16, 16); // 16x16 bloc(N + 15) / 16);
matrixMultiply<<<numBlocks, threadsPerBlock>>>(d_A, d_B, d_C, N);
cudaDeviceSynchronize();
matrixMultiply<<<numBlocks, threadsPerBlock>>>(d_C, d_D, d_E, N);
cudaDeviceSynchronize();
cudaMemcpy(h_E, d_E, size, cudaMemcpyDeviceToHost);

printf("Result matrix E (a few elements):\n");
for (int i = 0; i < 5; i++) {
    printf("%f ", h_E[i]);
}
printf("\n");
free(h_A);
free(h_B);
free(h_C);
free(h_D);
free(h_E);
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);
cudaFree(d_D);
cudaFree(d_E);

return 0;

```

Result matrix E (a few elements):

256.000000 320.000000 128.000000 512.000000 384.000000